

CS 181 Spring 2017 Section 3 Notes

(Bayesian Linear Regression, Classification)

1 Bayesian Linear Regression

1.1 Parameter Distributions

Let $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, $\mathbf{x}_i \in \mathbb{R}^m$, $y_i \in \mathbb{R}$. Consider the generative model

$$y_i \sim \mathcal{N}(\mathbf{w}^\top \mathbf{x}_i, \beta^{-1}) \quad (1)$$

The likelihood of the data has the form:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \mathcal{N}(\mathbf{y}|\mathbf{X}\mathbf{w}, \beta^{-1}\mathbf{I}) \quad (2)$$

Put conjugate prior on the weights as (assume precision β^{-1} known):

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_0, \mathbf{S}_0) \quad (3)$$

We want a posterior distribution on \mathbf{w} by using Bayes's Theorem, which states that:

$$p(\mathbf{w}|D) \propto p(D|\mathbf{w})p(\mathbf{w}) \quad (4)$$

It turns out (see practice question) that our posterior after n examples is also Gaussian:

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_n, \mathbf{S}_n) \quad (5)$$

where

$$\mathbf{S}_n = \left(\mathbf{S}_0^{-1} + \beta \mathbf{X}^\top \mathbf{X}\right)^{-1} \quad (6)$$

$$\mathbf{m}_n = \mathbf{S}_n(\mathbf{S}_0^{-1}\mathbf{m}_0 + \beta \mathbf{X}^\top \mathbf{y}) \quad (7)$$

1.2 Posterior Predictive Distributions

We don't just want to find the value of \mathbf{w} — we want to be able to predict y for new values of the input data. Let \mathbf{x} denote one such new data point.

For now, we assume that $p(\mathbf{w}) \sim \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I})$. Marginalizing over \mathbf{w} , we have that

$$p(y|\mathbf{x}, D) = \int_{\mathbf{w}} p(y|\mathbf{x}, \mathbf{w})p(\mathbf{w}|D)d\mathbf{w} \quad (8)$$

$$= \int_{\mathbf{w}} \mathcal{N}(y|\mathbf{w}^\top \mathbf{x}, \beta^{-1})\mathcal{N}(\mathbf{w}|\mathbf{m}_n, \mathbf{S}_n)d\mathbf{w} \quad (9)$$

Since each of the terms on the right hand side follows a normal distribution, we can use some math (see lecture 5, slide 33) to find that

$$p(y|\mathbf{x}, D) = \mathcal{N}(y|\mathbf{x}^\top \mathbf{m}_n, \mathbf{x}^\top \mathbf{S}_n \mathbf{x} + \beta^{-1}) \quad (10)$$

2 Practice Questions

1. Posterior Weight Distribution By Completing the Square (Bishop 3.7)

We know from (3.10) in Bishop that the likelihood can be written as

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{i=1}^n \mathcal{N}(y_i | \mathbf{w}^\top \mathbf{x}_i, \beta^{-1}) \\ \propto \exp \left(-\frac{\beta}{2} (\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w}) \right)$$

where precision $\beta = \frac{1}{\sigma^2}$ and in the second line above we have ignored the Gaussian normalization constants. By completing the square, show that with a prior distribution on \mathbf{w} given by $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_0, \mathbf{S}_0)$, the posterior distribution $p(\mathbf{w}|D)$ is given by

$$p(\mathbf{w}|D) = \mathcal{N}(\mathbf{w}|\mathbf{m}_n, \mathbf{S}_n)$$

where

$$\mathbf{m}_n = \mathbf{S}_n(\mathbf{S}_0^{-1}\mathbf{m}_0 + \beta\mathbf{X}^\top\mathbf{y}) \\ \mathbf{S}_n = \left(\mathbf{S}_0^{-1} + \beta\mathbf{X}^\top\mathbf{X} \right)^{-1}$$

2. Bayesian Updates in Linear Regression (Bishop 3.8)

Suppose we have the standard Bayesian linear regression model and we have already observed n data points, so the posterior distribution is

$$p(\mathbf{w}|D) = \mathcal{N}(\mathbf{w}|\mathbf{m}_n, \mathbf{S}_n)$$

where

$$\mathbf{m}_n = \mathbf{S}_n(\mathbf{S}_0^{-1}\mathbf{m}_0 + \beta\mathbf{X}^\top\mathbf{y})$$

$$\mathbf{S}_n = \left(\mathbf{S}_0^{-1} + \beta\mathbf{X}^\top\mathbf{X}\right)^{-1}$$

Suppose we observe a new data point $(\mathbf{x}_{n+1}, y_{n+1})$. Show that the resulting posterior distribution is of the same form with \mathbf{m}_{n+1} and \mathbf{S}_{n+1} .

3 Classification

The goal in classification is to take an input vector \mathbf{x} and assign it to one of c discrete classes, C_k , where $k = 1, \dots, c$. The input space is thus divided into **decision regions** whose boundaries are called **decision boundaries or surfaces**.

3.1 Binary Linear Classification

A discriminant function is one that directly assigns each vector \mathbf{x} to a specific class. We first assume two classes, i.e. our responses are binary and $c = 2$. Linear classification seeks to divide the 2 classes by a linear separator in the feature space— if $m = 2$ the separator is a line; if $m = 3$ the separator is a plane; for general m the separator is a $(m - 1)$ -dimensional hyperplane.

The simplest representation of a linear discriminant function is obtained by taking a linear function of the input vector as such:

$$h(\mathbf{x}; \mathbf{w}, w_0) = \mathbf{w}^\top \mathbf{x} + w_0$$

The corresponding decision boundary is defined by the relation $h(\mathbf{x}; \mathbf{w}, w_0) = 0$, which corresponds to the $(m - 1)$ -dimensional hyperplane within the d -dimensional input space.

The corresponding classifier will predict $\hat{y} = 1$ if $h(\mathbf{x}; \mathbf{w}, w_0) > 0$, and predict $\hat{y} = -1$ otherwise.

Weight vector \mathbf{w} is orthogonal to every vector lying within the decision surface, and so \mathbf{w} determines the orientation of the decision boundary. Mathematically, consider a vector $\mathbf{x}_1 - \mathbf{x}_2$ that lies on the boundary. We have:

$$\mathbf{w}^\top (\mathbf{x}_1 - \mathbf{x}_2) = -w_0 - (-w_0) = 0$$

From this, we see that \mathbf{w} is orthogonal to a vector that lies on the decision boundary (since the inner product is zero).

Moreover, w_0 (called the bias or threshold), determines the location of the decision boundary. In particular, we can ask for the normalized distance from a point \mathbf{x} on the boundary to the origin. By definition, the normalized distance from the origin to a point \mathbf{x} on the boundary is $\frac{\mathbf{w}^\top \mathbf{x}}{\|\mathbf{w}\|}$, where we consider the inner product between \mathbf{w} (the orthogonal vector to the plane) and the point \mathbf{x} . We divide by the norm $\|\mathbf{w}\|$, where $\|\mathbf{w}\|^2 = \mathbf{w}^\top \mathbf{w}$, to get the normalized distance. Noting that $h(\mathbf{x}; \mathbf{w}, w_0) = \mathbf{w}^\top \mathbf{x} + w_0 = 0$, since it is on the decision boundary, we substitute $\mathbf{w}^\top \mathbf{x} = -w_0$ to get

$$\frac{\mathbf{w}^\top \mathbf{x}}{\|\mathbf{w}\|} = -\frac{w_0}{\|\mathbf{w}\|},$$

and see that w_0 together with $\|\mathbf{w}\|$ define the location of the decision boundary. The input space can also be transformed through a basis technique, so that

$$h(\mathbf{x}; \mathbf{w}, w_0) = \mathbf{w}^\top \phi(\mathbf{x}) + w_0,$$

and this can help with linear separability.

3.2 Perceptron Algorithm

An important way to train a linear discriminant model is via the perceptron algorithm.

This works for a two-class model. Rather than a 0/1 error function (or sum-squared error), the perceptron algorithm adopts an alternative error function known as the rectified linear activation (ReLU), given by:

$$f_{relu}(z) = \begin{cases} z & z > 0 \\ 0 & o.w. \end{cases} = \max\{0, z\}$$

For a binary classification problem with classes 1 and -1, note that if $h(\mathbf{x}; \mathbf{w}, w_0) > 0$ (and thus $\hat{y} = 1$), then there is a classification error if the correct label is -1. Similarly (and ignoring when a point is right on the boundary) if $h(\mathbf{x}; \mathbf{w}, w_0) < 0$ (and thus $\hat{y} = -1$), then there is a classification error if the correct label is 1. For this reason, when the product $-h(\mathbf{x}_i; \mathbf{w}, w_0)y_i < 0$ then there is a classification error.

The perceptron loss function is defined as:

$$\begin{aligned} \mathcal{L}(\mathbf{w}) &= \sum_{i=1}^n f_{relu}(-h(\mathbf{x}_i; \mathbf{w}, w_0)y_i) \\ &= - \sum_{i=1: y_i \neq \hat{y}_i}^n (\mathbf{w}^\top \mathbf{x}_i + w_0)y_i \end{aligned}$$

The first term takes the sum over all training examples of the ReLU function applied to $-h(\mathbf{x}_i; \mathbf{w}, w_0)y_i$. In particular, when there is a misclassified example, then this value is positive and it counts as a loss. Equivalently, we can simply write this as the negated sum over all misclassified examples of $h(\mathbf{x}_i; \mathbf{w}, w_0)y_i$.

This loss function has a gradient that is easier to work with than if we had used a 0/1 error function, and we can now apply stochastic gradient descent. This keeps doing a gradient update on weights for one of the currently misclassified examples. The change in weight vector from step τ to $\tau + 1$ is given by the following iteration on an incorrect example:


$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \frac{\partial}{\partial \mathbf{w}} \mathcal{L}^{(i)}(\mathbf{w}) = \mathbf{w}^{(\tau)} + \eta y_i \mathbf{x}_i,$$

where η is the learning rate parameter. Note that as the weight vector evolves during training, the set of examples that are misclassified will also change.

3.3 Example of Perceptron Algorithm

Let the initial values be $\eta = 0.2^\dagger$, $\mathbf{w} = \begin{pmatrix} w_1 \\ w_2 \end{pmatrix} = \begin{pmatrix} 1 \\ 0.5 \end{pmatrix}$, $w_0 = 0$. The data $\{\mathbf{x}_i\}_{i=1}^n$, $\mathbf{x}_i \in \mathbf{R}^2$ and initial separation boundary are illustrated below. We choose the circles to belong to class 1 where ($y_i = 1$) and the triangles to belong to class 2 (with $y_i = -1$).

[†]For the perceptron, $\eta = 1$ almost all the time, but we set it differently for the sake of the exercise



Left as an exercise for the students.